

```

1 !pip install pandas numpy scikit-learn nltk textblob gensim spacy transformers matplotlib
2 !python -m spacy download en_core_web_sm
3 !python -m nltk.download punkt wordnet omw-1.4 stopwords

```

[Show hidden output](#)

```

1
2
3
4
5 import pandas as pd
6 import numpy as np
7 import re
8 import nltk
9 import spacy
10 from textblob import TextBlob
11 from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import accuracy_score, classification_report
14 import matplotlib.pyplot as plt
15 import seaborn as sns
16
17 # Deep Learning
18 import torch
19 from torch import nn
20 from torch.utils.data import DataLoader, TensorDataset
21 from transformers import BertTokenizer, BertForSequenceClassification, pipeline
22 from gensim.models import Word2Vec, FastText
23 from gensim.utils import simple_preprocess
24 import warnings
25 warnings.filterwarnings("ignore")
26
27 # NLTK & SpaCy
28 nltk.download('punkt')
29 nltk.download('punkt_tab')
30 nltk.download('stopwords')
31 nltk.download('wordnet')
32 nlp = spacy.load("en_core_web_sm")
33
34 # Set style
35 plt.style.use('seaborn-v0_8')
36 sns.set_palette("husl")

```

[Show hidden output](#)

## ▼ 1.Text Pre-processing on file.txt

```

1 # a. Import necessary libraries
2 import nltk
3 from nltk.tokenize import word_tokenize, sent_tokenize
4 from nltk.corpus import stopwords
5 from nltk.stem import PorterStemmer, WordNetLemmatizer
6 from textblob import TextBlob
7 import spacy
8
9 # Auto-create file.txt with improper text
10 file_content = """
11 Natural language processing (NLP) is a sub-field of artificial intelligence, concerned with
12
13 NLP is used in many applications such as: machine translation, sentiment analysis, speech
14
15 Some common challenges include: ambiguity in language, sarcasm, slang, and misspellings. For
16
17 Here are some sample sentences with errors:
18 - "Ths sentnce has spelng erors."
19 - "I luv programming in pyhton!!!"

```

```

20 - "Artificial intelligence is the future...""
21
22 NLP techniques help clean and normalize such texts. Tokenization breaks text into words, s-
23
24 Named entity recognition (NER) identifies names of people, organizations, locations. For e;
25
26 Finally, sentence boundary detection splits text into sentences. This is important for sum-
27 """
28
29 with open("file.txt", "w", encoding="utf-8") as f:
30     f.write(file_content.strip())
31
32 print("file.txt created successfully!")
33 nltk.download('punkt_tab')
34 nltk.download('averaged_perceptron_tagger_eng')
35 nltk.download('punkt')
36 nltk.download('stopwords')
37 nltk.download('wordnet')
38 nltk.download('averaged_perceptron_tagger')
39
40 nlp = spacy.load("en_core_web_sm")
41
42 # b. Load the text corpus
43 with open("file.txt", "r", encoding="utf-8") as f:
44     text = f.read()
45
46 print("Original text length:", len(text))
47 print(text[:500], "\n")
48
49 # c. Tokenization
50 tokens = word_tokenize(text.lower())
51 print("First 30 tokens:", tokens[:30])
52
53 # d. Spelling correction
54 corrected_tokens = []
55 for token in tokens:
56     corrected = str(TextBlob(token).correct())
57     corrected_tokens.append(corrected)
58
59 corrected_text = " ".join(corrected_tokens)
60 print("\nFirst 10 corrected tokens:", corrected_tokens[:10])
61 print("\nCorrected text (first 200 chars):", corrected_text[:200])
62
63 # e. POS tagging on corrected tokens
64 pos_tags = nltk.pos_tag(corrected_tokens)
65 print("\nPOS tags (first 10):", pos_tags[:10])
66
67 # f. Remove stop words
68 stop_words = set(stopwords.words("english"))
69 filtered_tokens = [t for t in corrected_tokens if t not in stop_words and t.isalpha()]
70 print("\nFirst 20 tokens after stop-word removal:", filtered_tokens[:20])
71
72 # g. Stemming & Lemmatization
73 stemmer = PorterStemmer()
74 lemmatizer = WordNetLemmatizer()
75
76 stemmed = [stemmer.stem(t) for t in filtered_tokens]
77 lemmatized = [lemmatizer.lemmatize(t) for t in filtered_tokens]
78
79 print("\nStemmed (first 20):", stemmed[:20])
80 print("Lemmatized (first 20):", lemmatized[:20])
81
82 # h. Sentence boundaries
83 sentences = sent_tokenize(text)
84 print(f"\nTotal number of sentences: {len(sentences)}")

```

```

file.txt created successfully!
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!

```

```
[nltk_data] Downloading package averaged_perceptron_tagger_eng to
[nltk_data]     /root/nltk_data...
[nltk_data] Unzipping taggers/averaged_perceptron_tagger_eng.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]     /root/nltk_data...
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!
Original text length: 1253
Natural language processing (NLP) is a sub-field of artificial intelligence, concerned with the interactions between computer and humans. NLP is used in many applications such as: machine translation, sentiment analysis, speech recognition, chatbots and information retrieval. Some common challenges include: ambiguity in language, sarcasm, slang, and misspellings.

First 30 tokens: ['natural', 'language', 'processing', '(', 'nlp', ')', 'is', 'a', 'sub-field', 'of', 'artificial', 'intelligence', 'concerned', 'with', 'the', 'interactions', 'between', 'computer', 'and', 'humans', '']

First 10 corrected tokens: ['natural', 'language', 'processing', '(', 'nap', ')', 'is', 'a', 'sub-field', 'of']

Corrected text (first 200 chars): natural language processing ( nap ) is a sub-field of artificial intelligence , concerned with the interactions between computer and humans.

POS tags (first 10): [('natural', 'JJ'), ('language', 'NN'), ('processing', 'NN'), ('(', '(', ('nap', 'JJ'), ')', ')'), (')', ')', ('', '')]

First 20 tokens after stop-word removal: ['natural', 'language', 'processing', 'nap', 'artificial', 'intelligence', 'concerned', 'with', 'the', 'interactions', 'between', 'computer', 'and', 'humans', '']

Stemmed (first 20): ['natur', 'languag', 'process', 'nap', 'artifici', 'intellig', 'concern', 'interact', 'comput', 'human', '']

Lemmatized (first 20): ['natural', 'language', 'processing', 'nap', 'artificial', 'intelligence', 'concerned', 'interaction']

Total number of sentences: 14
```

## 2. Feature Extraction on 20newsgroups

```
1 # a. Import packages
2 from sklearn.datasets import fetch_20newsgroups
3 import pandas as pd
4 import re
5 from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
6
7 # b. Fetch dataset
8 data = fetch_20newsgroups(subset='all', remove=('headers', 'footers', 'quotes'))
9 df = pd.DataFrame({'text': data.data, 'target': data.target})
10
11 # c. Clean data
12 def clean_text(txt):
13     txt = txt.lower()
14     txt = re.sub(r'^[^\w\s]', '', txt)
15     txt = re.sub(r'\s+', ' ', txt).strip()
16     return txt
17
18 df['cleaned'] = df['text'].apply(clean_text)
19
20 # d. BoW model
21 bow_vectorizer = CountVectorizer(max_features=5000, stop_words='english')
22 bow_matrix = bow_vectorizer.fit_transform(df['cleaned'])
23 bow_df = pd.DataFrame(bow_matrix.toarray(), columns=bow_vectorizer.get_feature_names_out())
24
25 # e. TF-IDF model
26 tfidf_vectorizer = TfidfVectorizer(max_features=5000, stop_words='english')
27 tfidf_matrix = tfidf_vectorizer.fit_transform(df['cleaned'])
28 tfidf_df = pd.DataFrame(tfidf_matrix.toarray(), columns=tfidf_vectorizer.get_feature_names_out())
29
30 # f. Compare top 20 frequent words
31 bow_sums = bow_df.sum().sort_values(ascending=False).head(20)
32 tfidf_sums = tfidf_df.sum().sort_values(ascending=False).head(20)
33
34 print("BoW Top 20 words:\n", bow_sums)
35 print("\nTF-IDF Top 20 words:\n", tfidf_sums)
```

BoW Top 20 words:

dont	6433
like	6391

### 3. Amazon Musical Instruments Reviews

```
1
2 # Import necessary libraries
3 import pandas as pd
4 import numpy as np
5 import re
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8
9 from sklearn.model_selection import train_test_split
10 from sklearn.feature_extraction.text import TfidfVectorizer
11 from sklearn.metrics import (
12     classification_report, confusion_matrix, accuracy_score,
13     roc_curve, auc, mean_squared_error
14 )
15
16 # Classifiers
17 from sklearn.linear_model import LogisticRegression
18 from sklearn.naive_bayes import MultinomialNB
19 from sklearn.neighbors import KNeighborsClassifier
20 from sklearn.tree import DecisionTreeClassifier
21 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
22 import xgboost as xgb
23
24 # Regressors
25 from sklearn.linear_model import LinearRegression
26 from sklearn.tree import DecisionTreeRegressor
27 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
28 import xgboost as xgb
29
30 # Set style
31 sns.set_palette("husl")
32
33 # =====
```

```

34 # a. Data Input
35 # =====
36
37 data = {
38     "reviewText": [
39         "Not much to write about here, but it does exac...",
40         "The product does exactly as it should and is q...",
41         "The primary job of this device is to block the...",
42         "Nice windscreen protects my MXL mic and preven...",
43         "This pop filter is great. It looks and perform...",
44         "So good that I bought another one. Love the h...",
45         "I have used monster cables for years, and with...",
46         "I now use this cable to run from the output of...",
47         "Perfect for my Epiphone Sheraton II. Monster ...",
48         "Monster makes the best cables and a lifetime w..."
49     ],
50     "Overall": [5, 4, 3, 5, 4, 4, 5, 5, 3, 4]
51 }
52
53 df = pd.DataFrame(data)
54 print("Original Data:")
55 print(df)
56 print("\n" + "*60 + "\n")
57
58 # =====
59 # b. Clean reviewText
60 # =====
61
62 def clean_review(text):
63     text = text.lower()
64     text = re.sub(r'^[a-z\s]', '', text) # Remove punctuation & numbers
65     text = re.sub(r'\s+', ' ', text).strip() # Normalize whitespace
66     return text
67
68 df['cleaned_review'] = df['reviewText'].apply(clean_review)
69
70 print("Cleaned Reviews (first 3):")
71 for i in range(3):
72     print(f"{i+1}. {df['cleaned_review'].iloc[i]}")
73 print("\n" + "*60 + "\n")
74
75 # =====
76 # c. TF-IDF Feature Extraction
77 # =====
78
79 vectorizer = TfidfVectorizer(
80     max_features=100,
81     stop_words='english',
82     ngram_range=(1,2), # Include bigrams for richer features
83     min_df=1
84 )
85
86 X = vectorizer.fit_transform(df['cleaned_review'])
87 feature_names = vectorizer.get_feature_names_out()
88
89 print(f"TF-IDF Matrix Shape: {X.shape}")
90 print(f"Features (sample): {feature_names[:10]}")
91 print("\n" + "*60 + "\n")
92
93 # Target for classification: 1 = Positive (4-5), 0 = Negative (1-3)
94 y_class = (df['Overall'] >= 4).astype(int) # Critical Fix: Use 0 and 1
95 y_reg = df['Overall']
96
97 # Train-test split
98 X_train_c, X_test_c, y_train_c, y_test_c = train_test_split(
99     X, y_class, test_size=0.3, random_state=42, stratify=y_class
100 )
101 X_train_r, X_test_r, y_train_r, y_test_r = train_test_split(

```

```

102     X, y_reg, test_size=0.3, random_state=42
103 )
104
105 # =====
106 # d. Models Definition
107 # =====
108
109 classifiers = {
110     "Logistic Regression": LogisticRegression(max_iter=1000),
111     "Naive Bayes": MultinomialNB(),
112     "KNN": KNeighborsClassifier(n_neighbors=3),
113     "Decision Tree": DecisionTreeClassifier(random_state=42),
114     "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
115     "GBM": GradientBoostingClassifier(random_state=42),
116     "XGBoost": xgb.XGBClassifier(
117         use_label_encoder=False,
118         eval_metric='logloss',
119         random_state=42
120     )
121 }
122
123 regressors = {
124     "Linear Regression": LinearRegression(),
125     "Decision Tree Reg": DecisionTreeRegressor(random_state=42),
126     "Random Forest Reg": RandomForestRegressor(n_estimators=100, random_state=42),
127     "GBM Reg": GradientBoostingRegressor(random_state=42),
128     "XGB Reg": xgb.XGBRegressor(random_state=42)
129 }
130
131 # =====
132 # e. Evaluation Function (Classification)
133 # =====
134
135 def evaluate_classification(model, X_test, y_test, name):
136     y_pred = model.predict(X_test)
137     y_prob = model.predict_proba(X_test)[:, 1] if hasattr(model, "predict_proba") else None
138
139     print(f"\n{'='*20} {name} {'='*20}")
140     print(f"Accuracy: {accuracy_score(y_test, y_pred):.3f}")
141     print("\nClassification Report:")
142     print(classification_report(y_test, y_pred, target_names=['Negative', 'Positive']))
143
144     # Confusion Matrix
145     cm = confusion_matrix(y_test, y_pred)
146     plt.figure(figsize=(5,4))
147     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
148                 xticklabels=['Neg', 'Pos'], yticklabels=['Neg', 'Pos'])
149     plt.title(f"{name} - Confusion Matrix")
150     plt.xlabel("Predicted")
151     plt.ylabel("Actual")
152     plt.show()
153
154     # ROC Curve
155     if y_prob is not None:
156         fpr, tpr, _ = roc_curve(y_test, y_prob)
157         roc_auc = auc(fpr, tpr)
158         plt.figure(figsize=(5,4))
159         plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.3f})')
160         plt.plot([0,1], [0,1], 'k--')
161         plt.xlim([0,1]); plt.ylim([0,1.05])
162         plt.xlabel('False Positive Rate')
163         plt.ylabel('True Positive Rate')
164         plt.title(f"{name} - ROC Curve")
165         plt.legend()
166         plt.show()
167
168 # =====
169 # CLASSIFICATION RESULTS

```

```
170 # =====
171
172 print("\nCLASSIFICATION RESULTS\n" + "-"*60)
173 for name, clf in classifiers.items():
174     clf.fit(X_train_c, y_train_c)
175     evaluate_classification(clf, X_test_c, y_test_c, name)
176
177 # =====
178 # REGRESSION RESULTS
179 # =====
180
181 print("\n\nREGRESSION RESULTS\n" + "-"*60)
182
183 for name, reg in regressors.items():
184     reg.fit(X_train_r, y_train_r)
185     y_pred = reg.predict(X_test_r)
186     mse = mean_squared_error(y_test_r, y_pred)
187     print(f"{name:25} → MSE: {mse:.3f} | RMSE: {np.sqrt(mse):.3f}")
188
189 # Optional: Predict on first review
190 print("\n" + "="*60)
191 print("SAMPLE PREDICTION (First Review):")
192 sample = vectorizer.transform([df['cleaned_review'].iloc[0]])
193 for name, clf in classifiers.items():
194     pred = clf.predict(sample)[0]
195     prob = clf.predict_proba(sample)[0]
```



Original Data:

	reviewText	Overall
0	Not much to write about here, but it does exac...	5
1	The product does exactly as it should and is q...	4
2	The primary job of this device is to block the...	3
3	Nice windscreen protects my MXL mic and prevent...	5
4	This pop filter is great. It looks and perform...	4
5	So good that I bought another one. Love the h...	4
6	I have used monster cables for years, and with...	5
7	I now use this cable to run from the output of...	5
8	Perfect for my Epiphone Sheraton II. Monster ...	3
9	Monster makes the best cables and a lifetime w...	4

=====

Cleaned Reviews (first 3):

- 1. not much to write about here but it does exac
- 2. the product does exactly as it should and is q
- 3. the primary job of this device is to block the

=====

TF-IDF Matrix Shape: (10, 70)

Features (sample): ['best' 'best cables' 'block' 'bought' 'bought love' 'cable' 'cable run' 'cables' 'cables lifetime' 'cables years']

=====

#### CLASSIFICATION RESULTS

===== Logistic Regression =====

Accuracy: 0.667

Classification Report:


	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 for this class. This occurs when there are no true positives.
```


```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 for this class. This occurs when there are no true positives.
```


```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 for this class. This occurs when there are no true positives.
```

Logistic Regression - Confusion Matrix



Logistic Regression - ROC Curve







```
===== Naive Bayes =====
Accuracy: 0.667
```

```
Classification Report:
```

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
```



```
===== KNN =====
```



```
Accuracy: 0.667
```

```
Classification Report:
```

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
```

```
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
```




```
===== Decision Tree =====
Accuracy: 0.667
```


Classification Report:

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 (use `average != 'binary'` to ignore this warning)
/warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 (use `average != 'binary'` to ignore this warning)
/warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 (use `average != 'binary'` to ignore this warning)
/warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
```



Decision Tree - ROC Curve




```
===== Random Forest =====
Accuracy: 0.667
```


Classification Report:

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 (use `average != 'binary'` to ignore this warning)
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 (use `average != 'binary'` to ignore this warning)
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and set to 0.0 (use `average != 'binary'` to ignore this warning)
  _warn_prf(average, modifier, f"{{metric.capitalize()}} is", len(result))
```



Random Forest - ROC Curve



```
===== GBM =====
Accuracy: 0.667
```


Classification Report:				
	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
```


```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
```

GBM - Confusion Matrix



GBM - ROC Curve



```
===== XGBoost =====
```

```
Accuracy: 0.667
```

Classification Report:				
	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	1
Positive	0.67	1.00	0.80	2
accuracy			0.67	3
macro avg	0.33	0.50	0.40	3
weighted avg	0.44	0.67	0.53	3

```
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [15:17:15] WARNING: /workspace/src/learner.c: Parameters: { "use_label_encoder" } are not used.
```


```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined since all predicted values are zero!
```

XGBoost - Confusion Matrix



## 4. IMDB Sentiment Analysis (Kaggle)

- 1.50

- 1.25

```
1 from google.colab import files
2 import os
3
4 if not os.path.exists("/root/.kaggle"):
5     os.makedirs("/root/.kaggle", exist_ok=True)
6
7 print("Please upload your kaggle.json file:")
8 uploaded = files.upload() # <-- Upload kaggle.json
9 for fn in uploaded.keys():
10    os.rename(fn, "/root/.kaggle/kaggle.json")
11 os.chmod("/root/.kaggle/kaggle.json", 600)
12
13 # -----
14 # 3. Download the exact dataset you mentioned
15 # -----
16 !kaggle datasets download -d columbine/imdb-dataset-sentiment-analysis-in-csv-format --un:
17
18 # The dataset contains Train.csv and Test.csv
19 # We will use Test.csv as per your instruction
20 csv_path = "Test.csv" # <-- This is the file we need
21 print(f"\nDataset ready: {csv_path}")
22
23 # -----
24 # 4. Import libraries
25 # -----
26 import pandas as pd
27 import re
28 from textblob import TextBlob
29 import warnings
30 warnings.filterwarnings("ignore")
31
32 # -----
33 # 5. Load the CSV file
34 # -----
35 df = pd.read_csv(csv_path)
36 print(f"Loaded {len(df)} rows from {csv_path}")
37 print(df.head())
38
39 # -----
40 # 6. Fetch the text column
41 # -----
42 # The column is named 'text' in this dataset
43 texts = df['text'].astype(str).copy()
44
45 # -----
46 # 7. Extract and remove @handles
47 # -----
48 def remove_handles(t):
49     return re.sub(r'@\w+', '', t)
50
51 texts_clean = texts.apply(remove_handles)
52
53 # -----
54 # 8. Perform sentiment analysis using TextBlob
55 # -----
56 def get_sentiment(txt):
57     blob = TextBlob(txt)
58     polarity = blob.sentiment.polarity
59     if polarity > 0.1:
60         return 'positive'
61     elif polarity < -0.1:
62         return 'negative'
63     else:
64         return 'neutral'
65
```